



PROXIMA

Experimental evaluation of optimal schedulers based on partitioned proportionate fairness

Davide Compagnin, Enrico Mezzetti and Tullio Vardanega
University of Padua - Italy



27th EUROMICRO Conference on Real-Time Systems (ECRTS)
Lund, July 9th, 2015

*This project and the research leading to these results
has received funding from the European
Community's Seventh Framework Programme [FP7 /
2007-2013] under grant agreement 611085*

www.proxima-project.eu

Outline

- ❑ Motivation of our work
- ❑ Brief recall of RUN and QPS algorithms
- ❑ Implementation and evaluation
- ❑ Conclusions and future work

Introduction

RUN

Reduction to UNiprocessor
(RTSS-11)

QPS

Quasi-Partitioning Scheduling
(ECRTS-14)

Optimal multiprocessor scheduling

Based on partitioned proportionate-fairness

Designed to reduce # of preemptions and migrations

On periodic task-sets

Also on sporadic task-sets

Motivation

RUN

Implemented¹
on top of LITMUS[^]RT

Confirming
moderate run-time overhead
in between that of P-EDF and G-EDF

QPS



¹Compagnin, D.; Mezzetti, E.; Vardanega, T., "Putting RUN into Practice: Implementation and Evaluation," (ECRTS-14)

Recall of the algorithms /1

RUN

QPS

Off-line phase

Multiprocessor
scheduling
problem

decomposition

Uniprocessor
scheduling
problems

On-line phase

The multiprocessor schedule is “derived” from the corresponding uniprocessor schedule

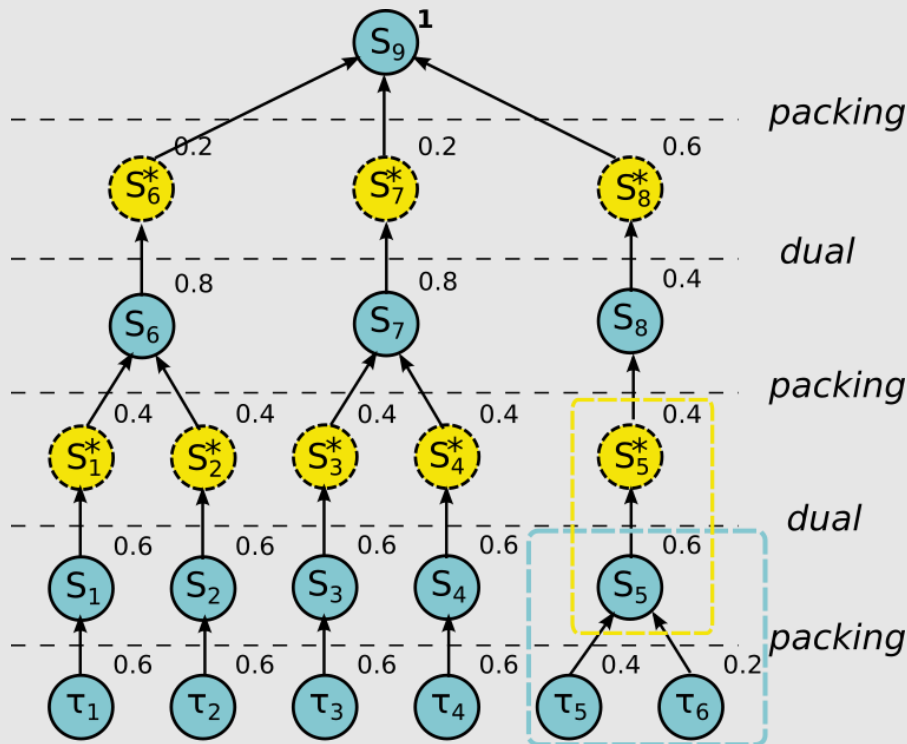
Recall of the algorithms /2

RUN

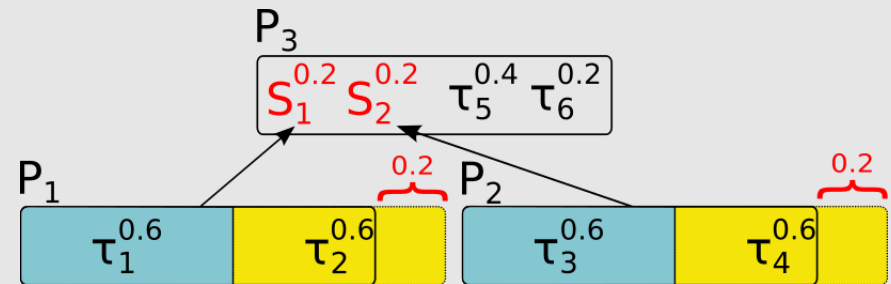
Off-line phase

QPS

Reduction tree



Processor hierarchy



Unitary processor capacity
can be exceeded

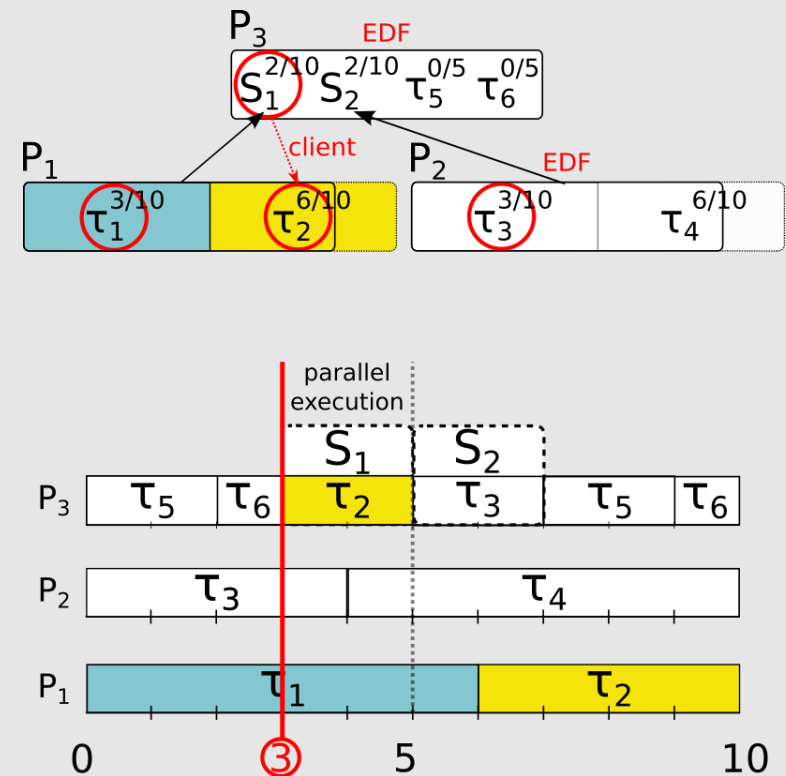
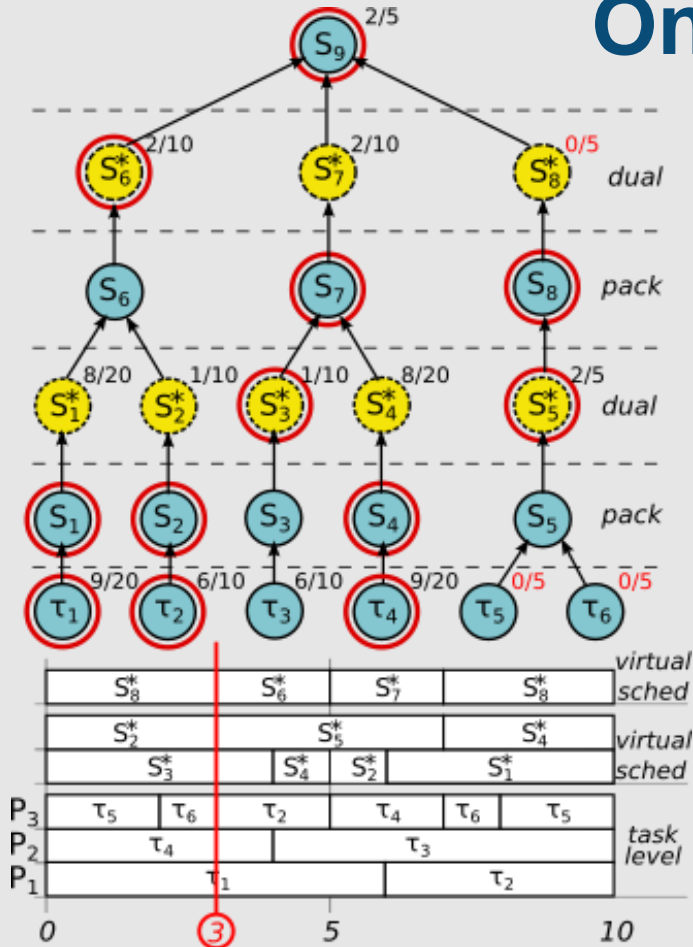
External servers
reserve capacity for exceeding
parts on a different processor

Recall of the algorithms /3

RUN

On-line phase

QPS

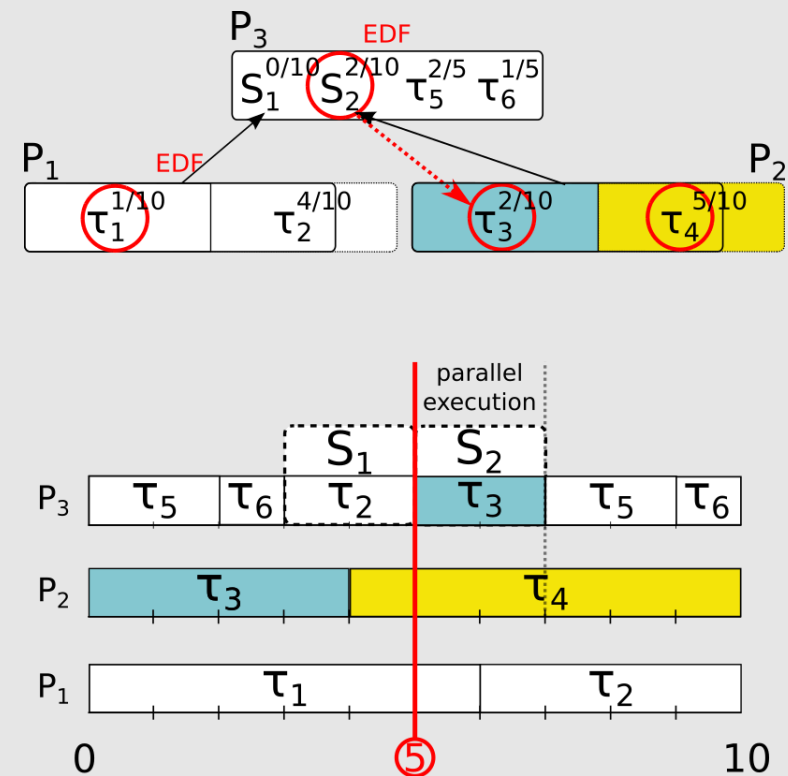
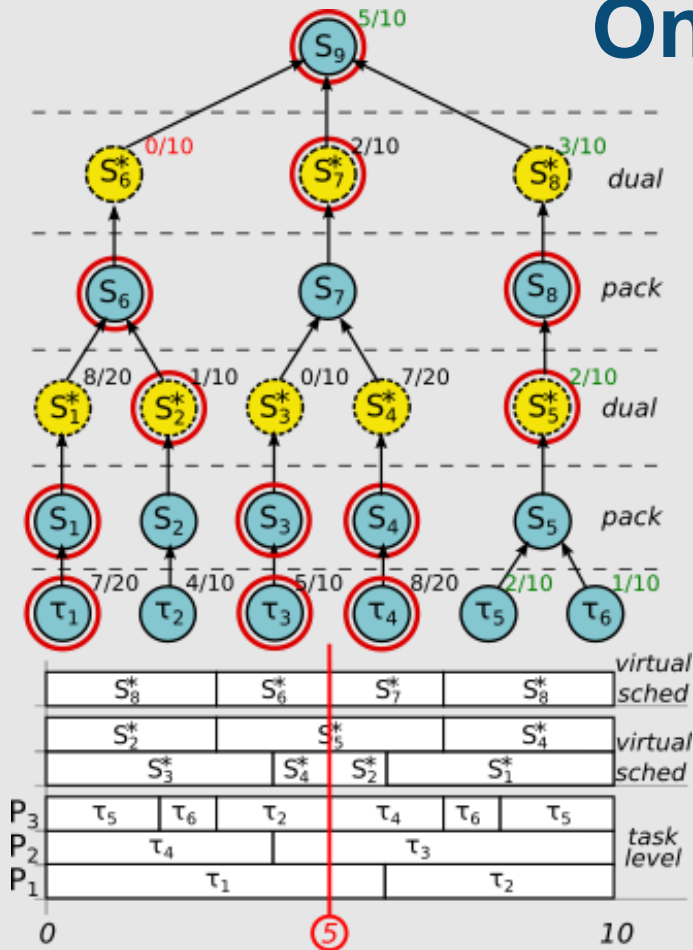


Recall of the algorithms /3

RUN

On-line phase

QPS



RUN

QPS

Notable differences

Global scheduling

- Virtual scheduling
- Compact tree representation
- CPUs are assigned to level-0 servers
- Timers trigger budget consumption events
- Node selection is performed
- Release queue and lock

Local scheduling

- With EDF

Local scheduling + Processor synchronization

- Uniform representation of tasks and servers
- Budgets consistently updated
- Timer triggers budget consumption events
- Per-hierarchy release queue and lock

Implementation /2

RUN

QPS

Notable differences

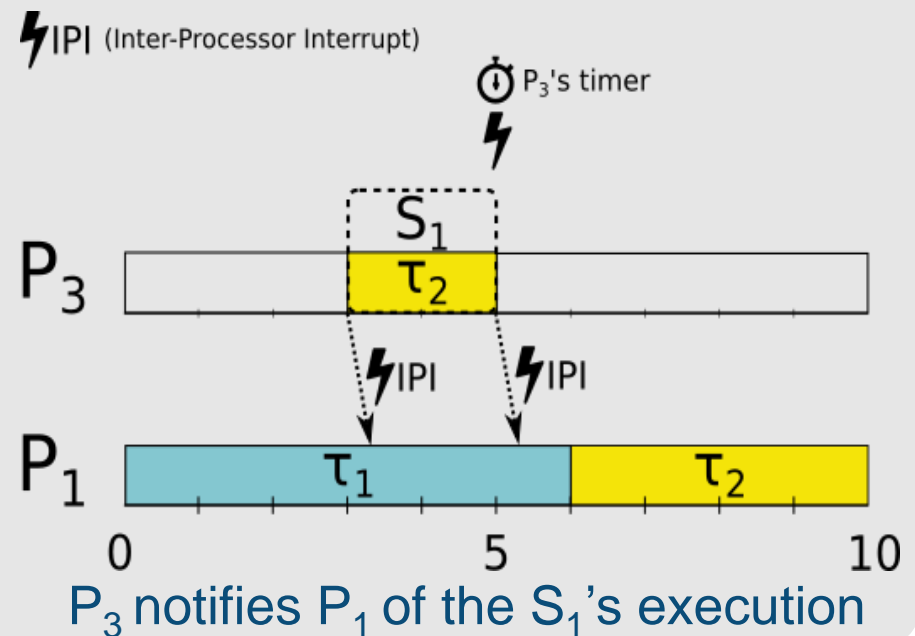
Global scheduling

- Virtual scheduling
- Compact tree representation
- CPUs are assigned to level-0 servers
- Timers trigger budget consumption events
- Node selection is performed
- Release queue and lock

Local scheduling

- With EDF

Local scheduling + Processor synchronization

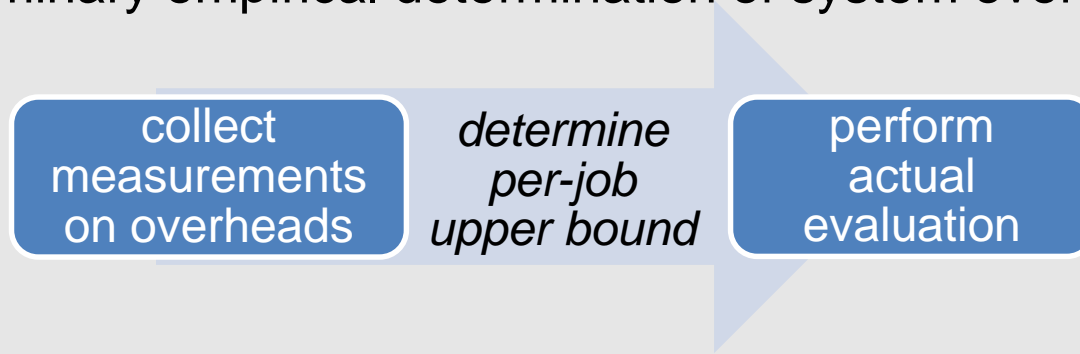


Evaluation

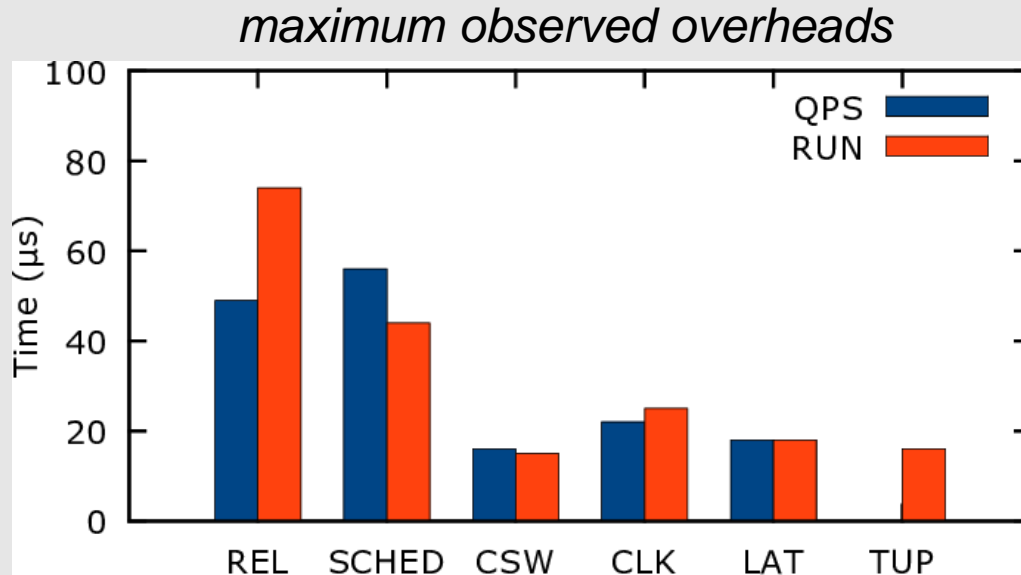
- ❑ Empirical evaluation instead of simulation
- ❑ Focus on scheduling interference
 - Cost of scheduling primitives
 - Incurred preemptions and migrations
- ❑ Evaluation limited to periodic task
 - *External servers* are always “active”
 - Sporadic activations would normally have lower utilization
 - Thus reducing the number of preemptions/migrations

Experimental setup

- ❑ LITMUS^{RT} on a 16-cores AMD Opteron 6370P
- ❑ Exhaustive measurements over the two algorithms
 - Thousand of automatically generated task sets
 - Harmonic and non-harmonic, with global utilization in 50%-100%
 - Stressing both the **off-line** and the **on-line** phases
- ❑ Two-step experimental process
 - Preliminary empirical determination of system overheads



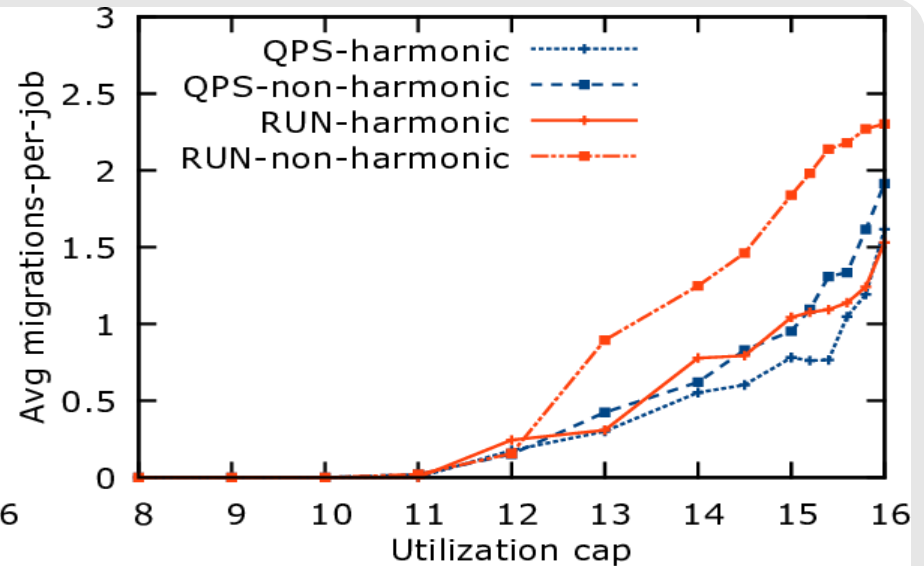
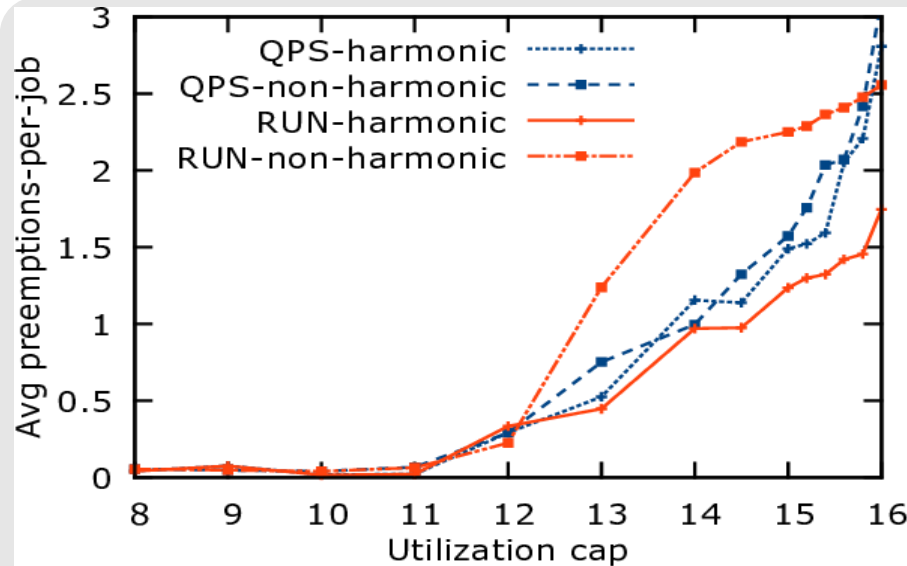
Primitive overheads and empirical bound



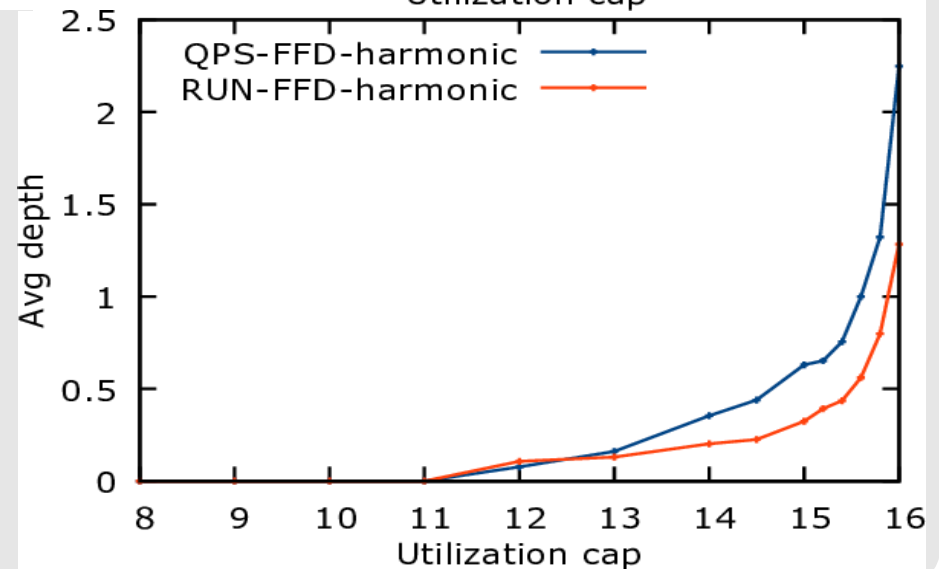
- ❑ Expectation was confirmed
 - QPS has lighter-weight scheduling primitives
 - And does not need Tree Update Operations (TUP)

- ❑ Empirical upper bound on the scheduling overhead
 - Based on theoretical bounds on the scheduling structures (RUN tree and QPS hierarchy)

Per-job scheduling interference

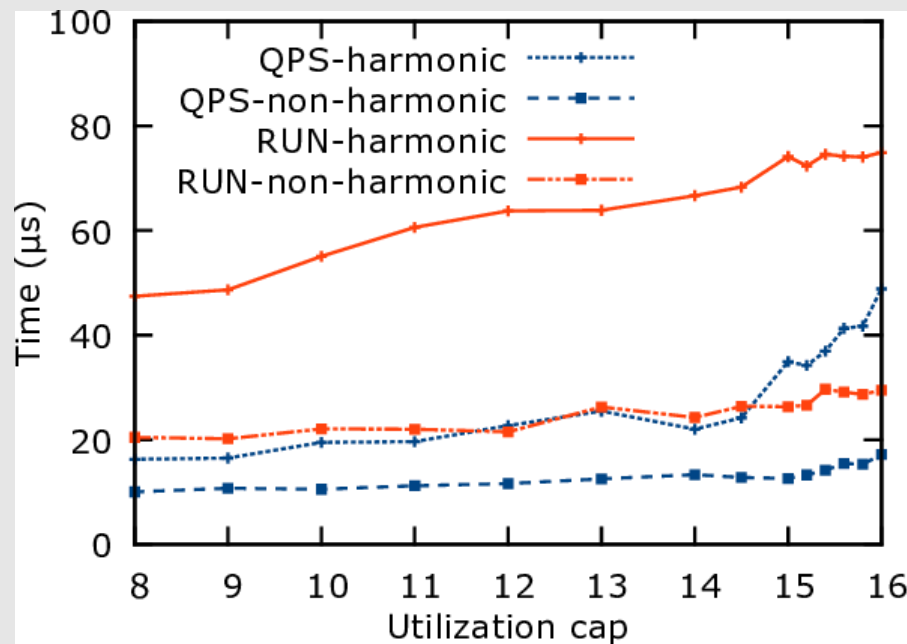


- Determined by preemptions and migrations
- In relation to reduction-tree and processor hierarchy depth

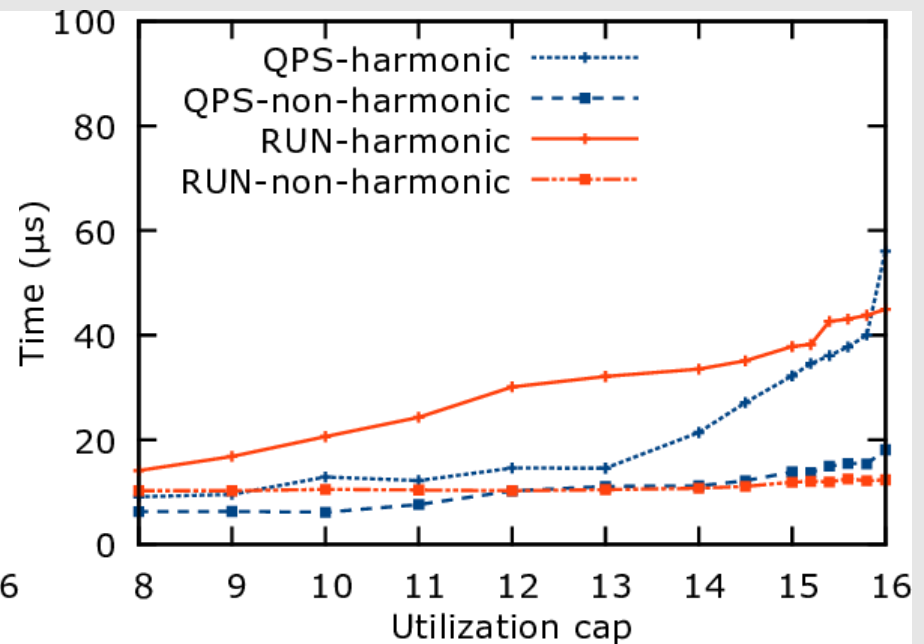


Scheduling primitives

max release



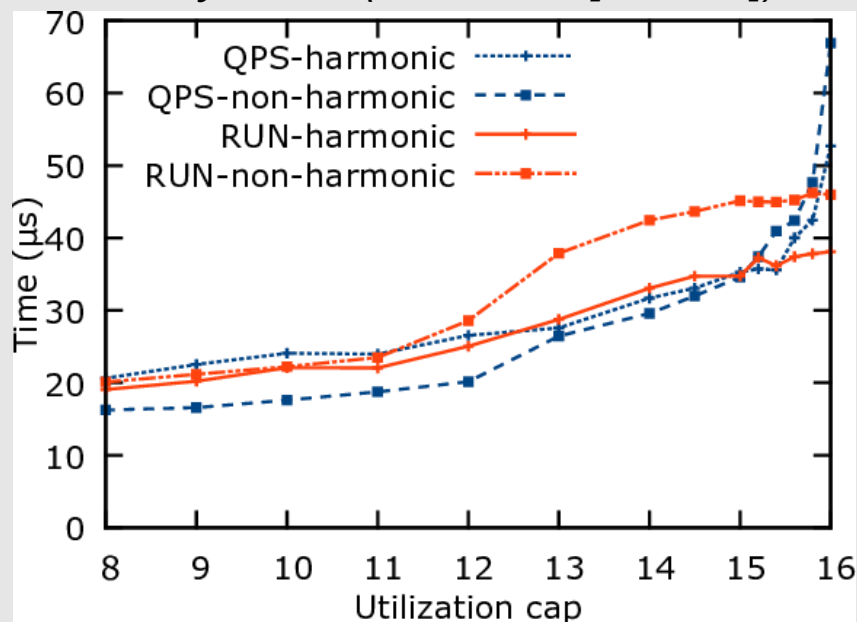
max schedule



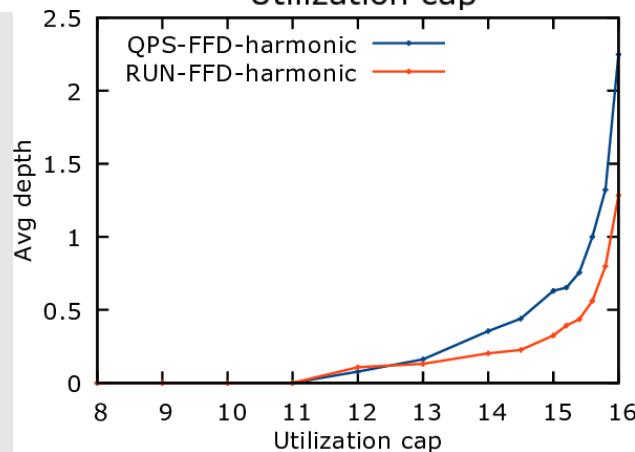
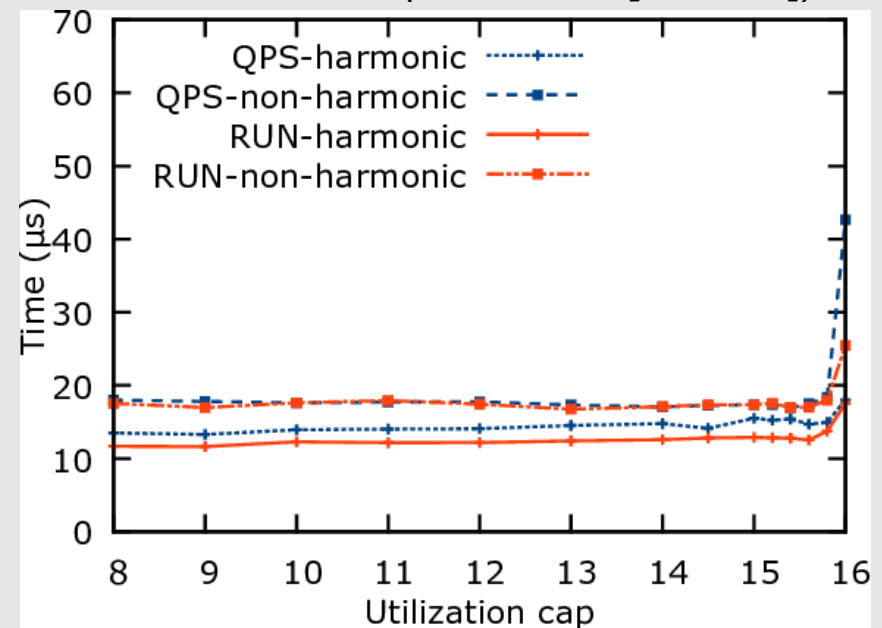
- ❑ Maximum observed cost of core scheduling primitives
 - *Release and Schedule*
 - Variation under increasing system utilization

Overall per-job overhead

heavy tasks (utilization [0.5;0.9])



medium tasks (utilization [0.1;0.5])



- QPS is more susceptible to packing than RUN
- Lighter-weight tasks ease the partitioning problem
 - And lead to less complex scheduling structures

Conclusions and future work

- ❑ QPS benefits from **partitioned scheduling**
 - Hence improves over RUN for cost of scheduling primitives
- ❑ ... but is more susceptible to the **off-line** phase
 - QPS's need for processor synchronization hits performance badly with higher processor hierarchies
- ❑ RUN exhibits an almost constant overhead
 - Induced by its global scheduling nature
 - Which in turn may penalize it at lower system utilization
- ❑ Future work
 - Mainly interested in evaluating how this class of algorithms may behave when the number of processing units increases
 - Considering also how different implementation may affect the algorithm scalability



PROXIMA

Experimental evaluation of optimal schedulers based on partitioned proportionate fairness

Davide Compagnin, Enrico Mezzetti and Tullio Vardanega
University of Padua - Italy



27th EUROMICRO Conference on Real-Time Systems (ECRTS)
Lund, July 9th, 2015

*This project and the research leading to these results
has received funding from the European
Community's Seventh Framework Programme [FP7 /
2007-2013] under grant agreement 611085*

www.proxima-project.eu